


```
# Importing the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

# getting the data from UC Irving machine learning repository
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/00468/online_shoppers_intention.csv"
df = pd.read_csv(url)
```

```
df.head()
```



	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration	BounceRa
0	0	0.0	0	0.0	1	0.000000	
1	0	0.0	0	0.0	2	64.000000	
2	0	0.0	0	0.0	1	0.000000	
3	0	0.0	0	0.0	2	2.666667	
4	0	0.0	0	0.0	10	627.500000	

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

Category	Column Name	Description	Values
Administrative	Administrative	Number of administrative pages visited	Integer
	Administrative_Duration	Time spent on administrative pages (seconds)	Float
Informational	Informational	Number of informational pages visited	Integer
	Informational_Duration	Time spent on informational pages (seconds)	Float
Product-Related	ProductRelated	Number of product pages visited	Integer
	ProductRelated_Duration	Time spent on product pages (seconds)	Float
User Engagement	BounceRates	Percentage of visitors who left immediately	Float (0–1)
	ExitRates	Percentage of pageviews that were last in session	Float (0–1)
	PageValues	Average value of pages viewed before purchase	Float (≥ 0)
Traffic Source	TrafficType	Source of traffic (20 categories)	Integer (1–20)
Visitor Type	VisitorType	Type of visitor	'New_Visitor', 'Returning_Visitor', 'Other'
Weekend	Weekend	Whether session occurred on weekend	Boolean (True/False)
Temporal	Month	Month of the year	Abbreviated (e.g., 'Feb', 'Nov')
Special Day	SpecialDay	Proximity to special shopping day	Float (0–1, where 1 = day of event)
Target Variable	Revenue	Whether purchase occurred	Boolean (True/False)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12330 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Administrative                        12330 non-null  int64
1   Administrative_Duration              12330 non-null  float64
2   Informational                        12330 non-null  int64
3   Informational_Duration               12330 non-null  float64
4   ProductRelated                      12330 non-null  int64
5   ProductRelated_Duration              12330 non-null  float64
6   BounceRates                         12330 non-null  float64
7   ExitRates                           12330 non-null  float64
8   PageValues                          12330 non-null  float64
9   SpecialDay                          12330 non-null  float64
10  Month                               12330 non-null  object
11  OperatingSystems                    12330 non-null  int64
12  Browser                             12330 non-null  int64
13  Region                             12330 non-null  int64
14  TrafficType                         12330 non-null  int64
15  VisitorType                         12330 non-null  object
16  Weekend                             12330 non-null  bool
17  Revenue                             12330 non-null  bool
dtypes: bool(2), float64(7), int64(7), object(2)
memory usage: 1.5+ MB
```

```
df.shape
```

```
(12330, 18)
```

```
df.isnull().sum()
```

```

0
Administrative      0
Administrative_Duration  0
Informational      0
Informational_Duration  0
ProductRelated     0
ProductRelated_Duration  0
BounceRates        0
ExitRates           0
PageValues          0
SpecialDay          0
Month              0
OperatingSystems    0
Browser             0
Region             0
TrafficType         0
VisitorType         0
Weekend             0
Revenue            0

```

```

duplicates = df[df.duplicated()]
duplicates.shape
print(duplicates)

```

```

Administrative  Administrative_Duration  Informational \
158            0                    0.0            0
159            0                    0.0            0
178            0                    0.0            0
418            0                    0.0            0
456            0                    0.0            0
...
11934          0                    0.0            0
11938          0                    0.0            0
12159          0                    0.0            0
12180          0                    0.0            0
12185          0                    0.0            0

Informational_Duration  ProductRelated  ProductRelated_Duration \
158                    0.0              1                    0.0
159                    0.0              1                    0.0
178                    0.0              1                    0.0
418                    0.0              1                    0.0
456                    0.0              1                    0.0
...
11934                  0.0              1                    0.0
11938                  0.0              1                    0.0
12159                  0.0              1                    0.0
12180                  0.0              1                    0.0
12185                  0.0              1                    0.0

BounceRates  ExitRates  PageValues  SpecialDay  Month  OperatingSystems \
158          0.2        0.2        0.0        0.0  Feb              1
159          0.2        0.2        0.0        0.0  Feb              3
178          0.2        0.2        0.0        0.0  Feb              3
418          0.2        0.2        0.0        0.0  Mar              1
456          0.2        0.2        0.0        0.0  Mar              2
...
11934        0.2        0.2        0.0        0.0  Dec              1
11938        0.2        0.2        0.0        0.0  Dec              1
12159        0.2        0.2        0.0        0.0  Dec              1
12180        0.2        0.2        0.0        0.0  Dec              1
12185        0.2        0.2        0.0        0.0  Dec              8

Browser  Region  TrafficType  VisitorType  Weekend  Revenue

```

```

158      1      1      3 Returning_Visitor False False
159      2      3      3 Returning_Visitor False False
178      2      3      3 Returning_Visitor False False
418      1      1      1 Returning_Visitor True False
456      2      4      1 Returning_Visitor False False
...      ...      ...      ...      ...      ...
11934    1      1      2      New_Visitor False False
11938    1      4      1 Returning_Visitor True False
12159    1      1      3 Returning_Visitor False False
12180    13     9     20 Returning_Visitor False False
12185    13     9     20      Other False False

```

[125 rows x 18 columns]

df.describe()

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration	Bou
count	12330.000000	12330.000000	12330.000000	12330.000000	12330.000000	12330.000000	12330.000000
mean	2.315166	80.818611	0.503569	34.472398	31.731468	1194.746220	12330.000000
std	3.321784	176.779107	1.270156	140.749294	44.475503	1913.669288	12330.000000
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	12330.000000
25%	0.000000	0.000000	0.000000	0.000000	7.000000	184.137500	12330.000000
50%	1.000000	7.500000	0.000000	0.000000	18.000000	598.936905	12330.000000
75%	4.000000	93.256250	0.000000	0.000000	38.000000	1464.157214	12330.000000
max	27.000000	3398.750000	24.000000	2549.375000	705.000000	63973.522230	12330.000000

df.dtypes

Administrative	int64
Administrative_Duration	float64
Informational	int64
Informational_Duration	float64
ProductRelated	int64
ProductRelated_Duration	float64
BounceRates	float64
ExitRates	float64
PageValues	float64
SpecialDay	float64
Month	object
OperatingSystems	int64
Browser	int64
Region	int64
TrafficType	int64
VisitorType	object
Weekend	bool
Revenue	bool

df.dtypes

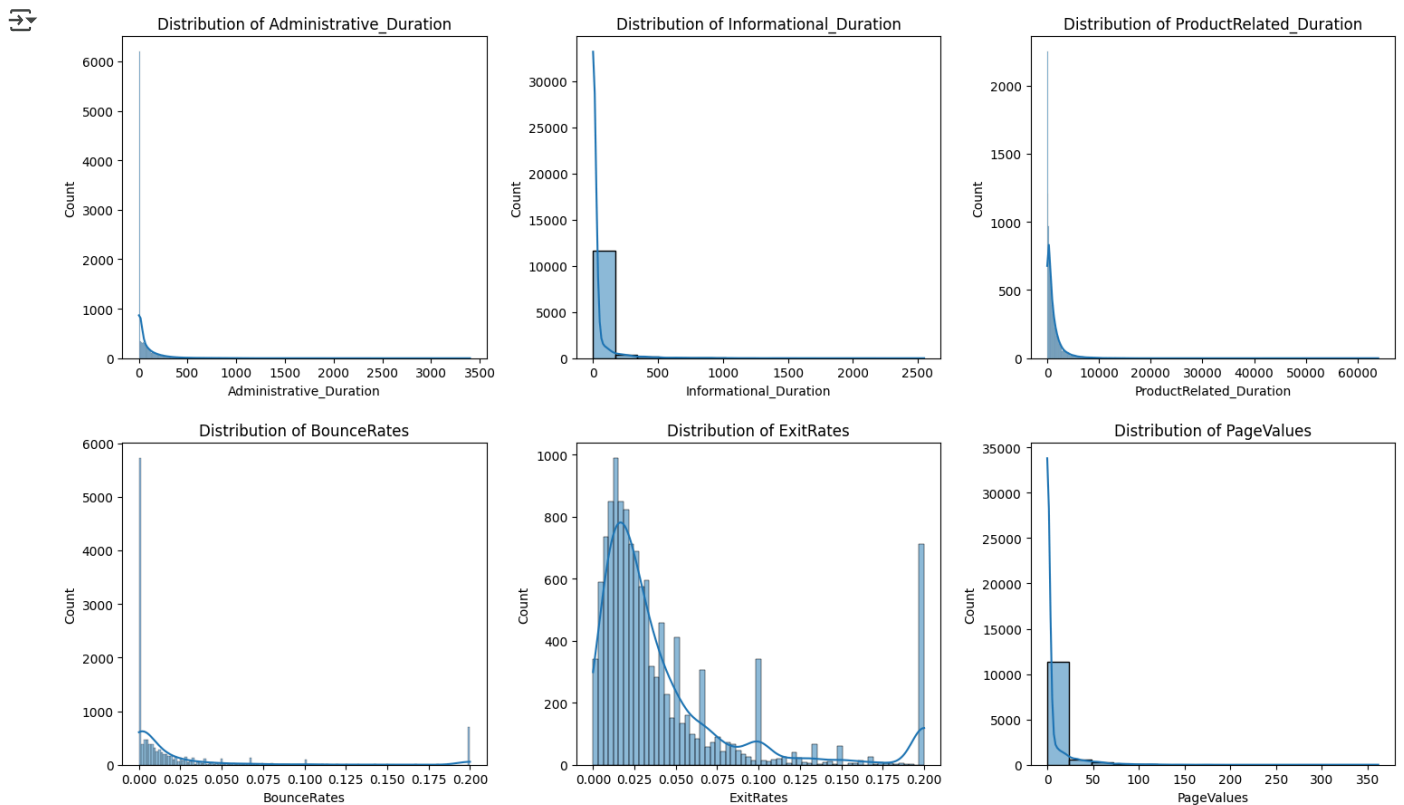
```

# Numerical features
numericalColumns = [
    'Administrative_Duration',
    'Informational_Duration',
    'ProductRelated_Duration',
    'BounceRates',
    'ExitRates',
    'PageValues'
]

plt.figure(figsize=(15, 9))
for i, col in enumerate(numericalColumns, 1):
    plt.subplot(2, 3, i)
    sns.histplot(df[col], kde=True)

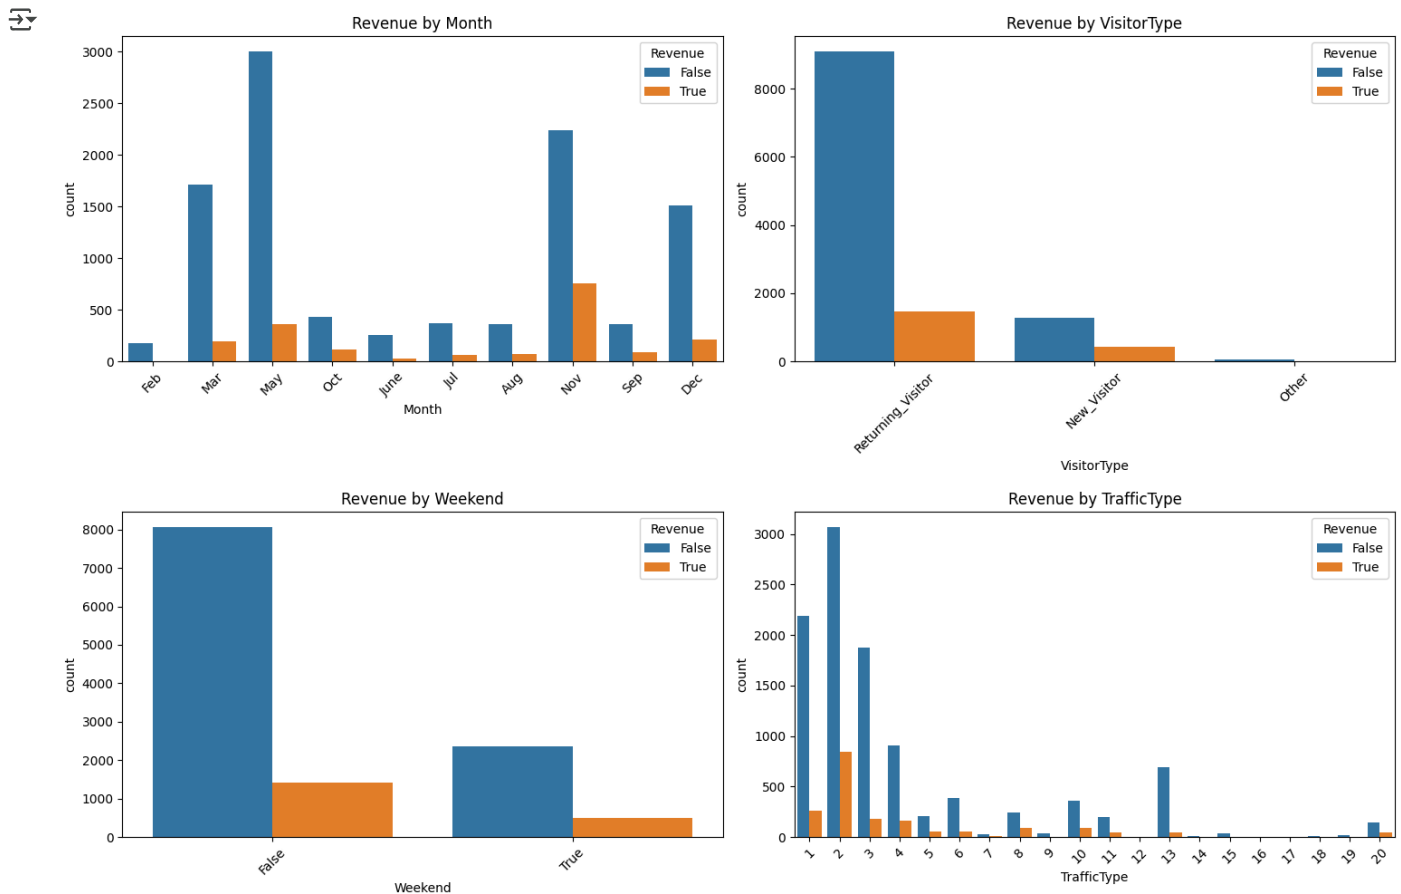
```

```
plt.title(f'Distribution of {col}')
plt.tight_layout(h_pad=2)
plt.show()
```



```
# Revenue distribution
categoricalColumns = [
    'Month',
    'VisitorType',
    'Weekend',
    'TrafficType'
]

plt.figure(figsize=(15, 10))
for i, col in enumerate(categoricalColumns, 1):
    plt.subplot(2, 2, i)
    sns.countplot(x=col, hue='Revenue', data=df)
    plt.title(f'Revenue by {col}')
    plt.xticks(rotation=45)
plt.tight_layout(h_pad=1.5)
plt.show()
```



```
# Correlation
corr = df[numericalColumns + ['Revenue']].corr()
corr
```


	Administrative_Duration	Informational_Duration	ProductRelated_Duration	BounceRates	ExitRates	PageValues
Administrative_Duration	1.000000	0.238031	0.355422	-0.144170	-0.205798	0.067608
Informational_Duration	0.238031	1.000000	0.347364	-0.074067	-0.105276	0.030861
ProductRelated_Duration	0.355422	0.347364	1.000000	-0.184541	-0.251984	0.052823
BounceRates	-0.144170	-0.074067	-0.184541	1.000000	0.913004	-0.119386
ExitRates	-0.205798	-0.105276	-0.251984	0.913004	1.000000	-0.174498
PageValues	0.067608	0.030861	0.052823	-0.119386	-0.174498	1.000000
Revenue	0.093587	0.070345	0.152373	-0.150673	-0.207071	0.492565

Next steps: [Generate code with corr](#) [View recommended plots](#) [New interactive sheet](#)

```
plt.figure(figsize=(12, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm', center=0, fmt='.2f')
plt.title('Feature Correlation Matrix')
plt.show()
```



Feature Correlation Matrix



```
# Feature engineering
```

```
df['Total_Duration'] = df['Administrative'] + df['Administrative_Duration'] + df['ProductRelated_Duration']
```

```
df['PageValue_to_ExitRatio'] = np.where(df['ExitRates'] > 0, df['PageValues'] / df['ExitRates'], 0)
```

```
df['BounceExit_Interaction'] = df['BounceRates'] * df['ExitRates']
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
df['Revenue'] = le.fit_transform(df['Revenue'])
```

```
df
```



	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration	Bour
0	0	0.0	0	0.0	1	0.000000	
1	0	0.0	0	0.0	2	64.000000	
2	0	0.0	0	0.0	1	0.000000	
3	0	0.0	0	0.0	2	2.666667	
4	0	0.0	0	0.0	10	627.500000	
...	
12325	3	145.0	0	0.0	53	1783.791667	
12326	0	0.0	0	0.0	5	465.750000	
12327	0	0.0	0	0.0	6	184.250000	
12328	4	75.0	0	0.0	15	346.000000	
12329	0	0.0	0	0.0	3	21.250000	

12330 rows × 21 columns

```

# Machine learning pipeline
from sklearn.model_selection import train_test_split, RandomizedSearchCV, StratifiedKFold
from sklearn.preprocessing import PowerTransformer, OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.metrics import classification_report, roc_auc_score, precision_recall_curve, average_precision_score, confusion_matrix, acc
import joblib
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier

x = df.drop('Revenue', axis=1)
y = df['Revenue']

# Splitting the data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42, stratify=y)

numericalFeatures = [
    'Administrative_Duration',
    'Informational_Duration',
    'ProductRelated_Duration',
    'BounceRates',
    'ExitRates',
    'PageValues',
    'Total_Duration',
    'PageValue_to_ExitRatio',
    'BounceExit_Interaction'
]

categoricalFeatures = [
    'Month',
    'VisitorType',
    'Weekend',
    'TrafficType',
]

preprocessor = ColumnTransformer(
    transformers=[
        ('num', Pipeline(steps=[
            ('power', PowerTransformer()),
            ('scaler', StandardScaler())
        ]), numericalFeatures),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categoricalFeatures)
    ]
)

smote = SMOTE(sampling_strategy=0.5, random_state=42)

models = {
    'LogisticRegression': {
        'model': LogisticRegression(max_iter=1000, class_weight='balanced'),
        'params': {
            'model__C': [0.1, 1, 10],
            'model__solver': ['lbfgs', 'liblinear']
        }
    },
    'NeuralNetwork': {
        'model': MLPClassifier(early_stopping=True),
        'params': {
            'model__hidden_layer_sizes': [(50,), (100,)],
            'model__alpha': [0.0001, 0.001],
            'model__learning_rate_init': [0.001, 0.01]
        }
    },
    'RandomForest': {
        'model': RandomForestClassifier(random_state=42),
        'params': {
            'model__n_estimators': [100, 200],
            'model__max_depth': [3, 5, 7]
        }
    },
    'GradientBoosting': {
        'model': GradientBoostingClassifier(random_state=42),
        'params': {
            'model__n_estimators': [100, 200],
            'model__learning_rate': [0.01, 0.1],
            'model__max_depth': [3, 5]
        }
    }
}

```

```

    }
}

results = {}

for name, config in models.items():
    print(f"\n====Training {name}====")

    # Create pipeline
    pipeline = Pipeline(steps=[
        ('preprocessor', preprocessor),
        ('smote', smote),
        ('model', config['model'])
    ])

    # Randomized Search
    search = RandomizedSearchCV(
        pipeline,
        config['params'],
        n_iter = 10,
        cv=StratifiedKFold(5),
        scoring='roc_auc',
        n_jobs=-1,
        random_state=42
    )

    search.fit(x_train, y_train)

    # Evaluate
    y_pred = search.best_estimator_.predict(x_test)
    y_proba = search.best_estimator_.predict_proba(x_test)[:,-1]

    # Store results
    results[name] = {
        'model': search.best_estimator_,
        'best_params': search.best_params_,
        'accuracy': accuracy_score(y_test, y_pred),
        'roc_auc': roc_auc_score(y_test, y_proba),
        'classification_report': classification_report(y_test, y_pred)
    }

    # Confusion matrix
    plt.figure()
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'{name} Confusion Matrix')
    plt.show()


    print(f"Best parameters: {search.best_params_}")
    print(classification_report(y_test, y_pred))

```




====Training LogisticRegression====

LogisticRegression Confusion Matrix




Best parameters: {'model__solver': 'liblinear', 'model__C': 0.1}

	precision	recall	f1-score	support
0	0.96	0.86	0.91	2084
1	0.52	0.81	0.63	382

accuracy			0.85	2466
macro avg	0.74	0.83	0.77	2466
weighted avg	0.89	0.85	0.86	2466

====Training NeuralNetwork====

NeuralNetwork Confusion Matrix




Best parameters: {'model__learning_rate_init': 0.001, 'model__hidden_layer_sizes': (100,), 'model__alpha': 0.0001}


	precision	recall	f1-score	support
0	0.95	0.90	0.92	2084
1	0.58	0.75	0.65	382

accuracy			0.88	2466
macro avg	0.76	0.82	0.79	2466
weighted avg	0.89	0.88	0.88	2466

====Training RandomForest====

RandomForest Confusion Matrix






Best parameters: {'model__n_estimators': 200, 'model__max_depth': 7}

	precision	recall	f1-score	support
0	0.95	0.90	0.93	2084
1	0.58	0.74	0.65	382
accuracy			0.88	2466
macro avg	0.77	0.82	0.79	2466
weighted avg	0.89	0.88	0.88	2466

====Training GradientBoosting====

GradientBoosting Confusion Matrix



Best parameters: {'model__n_estimators': 100, 'model__max_depth': 3, 'model__learning_rate': 0.1}

	precision	recall	f1-score	support
0	0.94	0.93	0.94	2084
1	0.64	0.70	0.67	382
accuracy			0.89	2466
macro avg	0.79	0.81	0.80	2466
weighted avg	0.90	0.89	0.89	2466

```
best_model = results['GradientBoosting']['model']  
best_model
```

